



Athena: Domain-Specific FAQ Chatbot with Knowledge Graph Integration

Traditional chatbots rely on predefined responses or general-purpose AI models, resulting in shallow answers lacking domain-specific depth. They struggle with complex relationships and contextual dependencies, making them unreliable for specialized use cases.

Our Proposed Approach

Dynamic Knowledge Graph

Information stored as interconnected entities and relationships, continuously updated in real-time from documents and data streams.

Neuro-Symbolic Fusion

Combines neural models (for understanding natural language) with symbolic logic (for consistent reasoning and enforcing rules).

Graph RAG

Intelligently traverses the graph to assemble context-rich, fact-grounded answers instead of vague summaries.

Explainability & Adaptability

Every answer traced to its source, ensuring trust.
Capable of ingesting diverse content and evolving with user feedback.

Athena is a reasoning-driven, domain-adaptive digital assistant that goes far beyond a traditional FAQ bot by constructing a dynamic knowledge graph and applying neuro-symbolic AI.

Key Features



Knowledge Ingestion

Reads FAQs, PDFs, websites, and converts raw text into structured triples.



Dynamic Knowledge Graph

Continuously updates, prunes, and enriches knowledge, enabling cross-domain reasoning.



Graph-RAG Reasoning

Traverses the graph to give context-rich answers, moving beyond keyword search.



Explainable Outputs

Every answer includes a reasoning trail with source links for verification.




Adaptive Learning

Learns from user corrections and refines knowledge continuously.



Reliability & Security

Ensures high availability, data privacy, and secure knowledge management.



How It Works

Ingestion & Enrichment

Multi-source integration with neuro-symbolic processing: Neural "System 1" extracts entities and relationships; Symbolic "System 2" applies logic and rules.

Dynamic Knowledge Graph

A living model of your domain that continuously updates in real-time, ensuring knowledge is never outdated.

Metta Language

Specialized language designed to build, manage, and reason about knowledge graphs, enabling the AI to learn and adapt.



GraphRAG

Hybrid retrieval process combining semantic search with graph traversal to generate comprehensive, accurate answers.

Key Challenges & Our Solutions

Challenge	Our Solution
Scalability: Knowledge graphs can grow to billions of facts, overwhelming traditional databases.	Distributed, Cloud-Native Architecture: Uses distributed graph databases and intelligent data partitioning to scale horizontally.
Data Integrity: Risk of creating a "knowledge swamp" of conflicting, outdated information.	Automated Governance: Enforces strict data quality with schema validation, ACID-compliant updates, and automated conflict resolution.
Trust & Bias: AI models can inherit and amplify harmful biases, making decisions opaque.	Explainability by Design: Every fact is traceable to its source, and the system can explain its reasoning path through the graph.
Computational Cost: Neuro-symbolic reasoning is expensive and inefficient on standard hardware.	Specialized Software & Hardware: Uses Metta language for efficient reasoning and helps define requirements for next-gen AI accelerators.

Athena vs. Current LLM-Based Systems (IS TABLE BETTER????)



Truth Model

Current: Generative
Plausibility - creates answers that sound correct.

Athena: Symbolic Verification
- constructs answers from proven facts.

Benefit: Unprecedented Trust & Reliability

Context

Current: Shallow Retrieval - finds semantically similar text passages.

Athena: Deep Relational Reasoning - understands hierarchies/causality.

Benefit: Actionable Insights, Not Just Summaries

Knowledge & Efficiency

Current: Static & Stale with high computational cost.

Athena: Dynamic & Real-Time with low query cost.

Benefit: Always Current & Massive Cost Savings

Business Model: QueryVerse



Core Model

Enterprise B2B SaaS with tiered pricing (Pilot, Professional, Enterprise) based on usage and number of knowledge bases.

Primary Value Driver

De-risking critical decisions by providing verifiable, accurate, and context-aware answers.

Target Customer

Organizations where the cost of being wrong is extremely high (legal, compliance, engineering, life sciences).

Competitive Moat

Hybrid KG + LLM architecture creates a self-improving knowledge asset that becomes more valuable over time.

Target Applications

1

Initial Beachhead (PoC)

Disaster Management & Emergency Response: Enabling first responders to get immediate, accurate answers from operational manuals during a crisis.

2

Legal & Compliance

Automating due diligence and accelerating case law research by transforming static legal documents into queryable evidence.

3

Pharmaceuticals & Life Sciences

Empowering researchers to analyze and cross-reference clinical trial data, patents, and scientific papers instantly.

4

High-Value Engineering

Providing engineers with precise, context-aware information from complex technical manuals and standards.

Ultimate Vision: Turn any organization's entire library of internal documentation into an interactive, trustworthy, and intelligent knowledge engine.



Future Scope: Project Athena

Phase 1: Immediate (Now → 2 Years)

The Proactive Expert

- Proactive assistance with real-time alerts
- Multi-modal understanding (text, tables, charts)
- Cross-domain synthesis connecting disparate knowledge graphs

Phase 3: Long-Term (5+ Years)

The AGI Foundation

- Autonomous knowledge ecosystem
- Core foundation for general intelligence

1

2

3

Phase 2: Mid-Term (2 → 5 Years)

The Analytical Partner

- Automated hypothesis generation
- Causal & predictive reasoning for "what-if" scenarios

Conclusion



Athena represents a paradigm shift in AI-powered knowledge systems by:

- Moving beyond traditional chatbots to deep reasoning with knowledge graphs
- Providing verifiable, trustworthy answers with full explainability
- Creating a self-improving knowledge asset that grows more valuable over time
- Offering a scalable foundation for mission-critical applications where accuracy matters

With QueryVerse's B2B SaaS model, we're positioned to transform how organizations access, understand, and leverage their knowledge assets.